

# RIPQ: Advanced Photo Caching on Flash for Facebook

Linpeng Tang (Princeton)

Qi Huang (Cornell & Facebook)


Wyatt Lloyd (USC & Facebook)

Sanjeev Kumar (Facebook)

Kai Li (Princeton)



PRINCETON  
UNIVERSITY




facebook®

# 2 Billion\* Photos Shared Daily



# Photo Serving Stack




# Photo Caches

**Close to users**  
Reduce backbone traffic

Co-located with backend  
Reduce backend IO

# Photo Serving Stack



# An Analysis of Facebook Photo Caching [Huang et al. SOSP'13]




Advanced caching  
algorithms help!

Segmented LRU-3:  
10% less backbone traffic

Greedy-Dual-Size-Frequency-3:  
23% fewer backend IOs

## Photo Serving Stack




# In Practice



FIFO was still used  
No known way to implement advanced algorithms efficiently

# Photo Serving Stack



# Theory



Advanced caching helps

- 23% fewer backend I/O
- 10% less backbone traffic

# Practice



Difficult to implement on flash:

I/O still used



**Restricted Insertion Priority Queue:**  
efficiently implement advanced  
caching algorithms on flash


# Outline

- Why are advanced caching algorithms difficult to implement on flash efficiently?
- How RIPQ solves this problem?
  - Why use priority queue?
  - How to efficiently implement one on flash?
- Evaluation
  - 10% less backbone traffic
  - 23% fewer backend IOs


# Outline

- Why are advanced caching algorithms difficult to implement on flash efficiently?
  - Write pattern of FIFO and LRU
- How RIPQ solves this problem?
  - Why use priority queue?
  - How to efficiently implement one on flash?
- Evaluation
  - 10% less backbone traffic
  - 23% fewer backend IOs


# FIFO Does Sequential Writes




# FIFO Does Sequential Writes



# FIFO Does Sequential Writes




# FIFO Does Sequential Writes




No random writes needed for FIFO

# LRU Needs Random Writes



Locations on flash  $\neq$  Locations in LRU queue

# LRU Needs Random Writes



Random writes needed to reuse space

# Why Care About Random Writes?

- **Write-heavy workload**
  - Long tail access pattern, moderate hit ratio
  - Each miss triggers a write to cache
- **Small random writes are harmful for flash**
  - e.g. Min et al. FAST'12
  - High write amplification
    - Low write throughput
    - Short device lifetime

# What write size do we need?


- Large writes
  - High write throughput at high utilization
  - 16~32MiB in Min et al. FAST'2012
- What's the trend since then?
  - Random writes tested for 3 modern devices
  - 128~512MiB needed now

100MiB+ writes needed for efficiency

# Outline

- Why are advanced caching algorithms difficult to implement on flash efficiently?
- How RIPQ solves this problem?
- Evaluation


# RIPQ Architecture (Restricted Insertion Priority Queue)



*Caching algorithms approximated as well*

*Efficient caching on flash*

# RIPQ Architecture (Restricted Insertion Priority Queue)




# Priority Queue API

- No single best caching policy
- Segmented LRU [Karedla'94]
  - Reduce both backend IO and backbone traffic
  - SLRU-3: best algorithm for Edge so far
- Greedy-Dual-Size-Frequency [Cherkasova'98]
  - Favor small objects
  - Further reduces backend IO
  - GDSF-3: best algorithm for Origin so far


# Segmented LRU

- Concatenation of K LRU caches




# Segmented LRU

- Concatenation of K LRU caches




# Segmented LRU

- Concatenation of K LRU caches



# Segmented LRU


- Concatenation of K LRU caches



# Greedy-Dual-Size-Frequency

- Favoring small objects

Cache space of GDSF-3



# Greedy-Dual-Size-Frequency

- Favoring small objects

Cache space of GDSF-3




Miss



# Greedy-Dual-Size-Frequency

- Favoring small objects

Cache space of GDSF-3




Miss



# Greedy-Dual-Size-Frequency


- Favoring small objects

Cache space of GDSF-3




- Write workload more random than LRU
- Operations similar to priority queue

# Relative Priority Queue for Advanced Caching Algorithms




Miss object:  $\text{insert}(x, p)$  

# Relative Priority Queue for Advanced Caching Algorithms



Hit object:  $\text{increase}(x, p')$


# Relative Priority Queue for Advanced Caching Algorithms



Implicit demotion on insert/increase:


- Object with lower priorities moves towards the tail

# Relative Priority Queue for Advanced Caching Algorithms




Relative priority queue captures the dynamics of many caching algorithms!

# RIPQ Design: Large Writes




- Need to buffer multiple writes (e.g., 100s) into block writes
- Once written, block is flushed to disk immediately!
- 256MiB block size, 16GiB memory allocation
  - Large caching capacity
  - High write throughput

# RIPQ Design: Restricted Insertion Points




- Exact priority queue
  - Insert to any block in the queue
- Each block needs a separate buffer
  - Whole flash space buffered in RAM!

# RIPQ Design: Restricted Insertion Points




**Solution: restricted insertion points**

# Section is Unit for Insertion



Each section has one insertion point


# Section is Unit for Insertion



## Insert procedure


- Find corresponding section
- Copy data into active block
- Updating section priority range

# Section is Unit for Insertion



Relative orders within one section not guaranteed!

# Trade-off in Section Size




Section size controls approximation error

- Sections ↑, approximation error ↓
- Sections ↓, RAM buffer ↓

# RIPQ Design: Lazy Update


Naïve approach: copy to the corresponding active block



Problem with naïve approach


- Data copying/duplication on flash

# RIPQ Design: Lazy Update




**Solution: use **virtual block** to track the updated location!**

# RIPQ Design: Lazy Update




**Solution: use **virtual block** to track the updated location!**


# Virtual Block Remembers Update Location



# Actual Update During Eviction



# Actual Update During Eviction



Always one copy of data on flash

# RIPQ Design

- Relative priority queue API
- RIPQ design points
  - Large writes
  - Restricted insertion points
  - Lazy update
  - Section merge/split
    - Balance section sizes and RAM buffer usage
- Static caching
  - Photos are static

# Outline

- Why are advanced caching algorithms difficult to implement on flash efficiently?
- How RIPQ solves this problem?
- Evaluation


# Evaluation Questions

- How much RAM buffer needed?
- How good is RIPQ's approximation?
- What's the throughput of RIPQ?


# Evaluation Approach

- Real-world Facebook workloads
  - Origin
  - Edge
- 670 GiB flash card
  - 256MiB block size
  - 90% utilization
- Baselines
  - FIFO
  - SIPQ: Single Insertion Priority Queue


# RIPQ Needs Small Number of Insertion Points



# RIPQ Needs Small Number of Insertion Points




# RIPQ Needs Small Number of Insertion Points




You don't need much RAM buffer (2GiB)!


# RIPQ Has High Fidelity



# RIPQ Has High Fidelity




# RIPQ Has High Fidelity




RIPQ achieves  $\leq 0.5\%$  difference for all algorithms

# RIPQ Has High Fidelity



+16% hit-ratio → 23% fewer backend IOs

# RIPQ Has High Throughput



**RIPQ throughput comparable to FIFO ( $\leq 10\%$  diff.)**

# Related Works

## RAM-based advanced caching

SLRU(Karedla'94), GDSF(Young'94, Cao'97, Cherkasova'01),  
SIZE(Abrams'96), LFU(Maffeis'93), LIRS (Jiang'02), ...

**RIPQ enables their use on flash**

## Flash-based caching solutions

Facebook FlashCache, Janus(Albrecht '13), Nitro(Li'13),  
OP-FCL(Oh'12), FlashTier(Saxena'12), Hec(Yang'13), ...

**RIPQ supports advanced algorithms**

## Flash performance

Stoica'09, Chen'09, Bouganim'09, Min'12, ...

**Trend continues for modern flash cards**

# RIPQ

- First framework for advanced caching on flash
  - Relative priority queue interface
  - Large writes
  - Restricted insertion points
  - Lazy update
  - Section merge/split
- Enables SLRU-3 & GDSF-3 for Facebook photos
  - 10% less backbone traffic
  - 23% fewer backend IOs